

Practical session 2: gradient descent

Optimization Techniques, UPF

April/May, 2023

In this assignment we will study the *gradient descent* algorithm, one of the simplest (and more general) function minimization methods. First we will consider a toy problem: the minimization of a function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$. The second part of the practice is on the application of the gradient descent method to remove the noise in an image, via the minimization of a denoising energy.

The day of the practice, you will be given some incomplete Python functions. Following the different parts of the practice you will complete them. You are not supposed to know the answers to all the exercises. We will guide you during the practice. For any doubts before and after the practice, you can contact us.

Pre-requisites: Before the practice, you should review the following topics:

- Gradient of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$.
- Level sets (or level lines) of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, and its geometrical relation to the gradient.
- Directional derivatives, and how to compute them using the gradient.

Deadline P101: Friday, May 5th at 23:55.

Deadline P102: Wednesday, May 10th at 23:55.

Grading: The evaluation is based on the report documenting your work (with figures), results, conclusions and the commented code. For example,

- The goal of the lab.
- Summary with your own words of the topic.
- Conclusions for each exercise.

Nneka Okolo: nnekamaureen.okolo@upf.edu

1 Minimization of a toy function

We will use the gradient descent method to compute the minima of the function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ given by,

$$f(x_1, x_2) = \frac{1}{1000} (x_1^4 + x_2^4 - 80x_1^2 - 60x_2^2 + 100x_1 + 50x_2 + 1)$$

Assignments

1. Complete the Python functions `toy_fun` and `toy_gradient`. These functions implement the function f and its gradient. Follow the comments provided in the code.
2. Complete the Python function `gradient_descent`. This function implements a gradient descent algorithm. We are going to implement it in a way in which we can use the same gradient descent function for this toy example and for the denoising energy of the next section. Follow the comments provided in the code.
3. Run the function `toy_main` with several time steps and several initial conditions and answer to the following questions:
 - How many local minima does the function have in the domain $[-10, 10] \times [-10, 10]$?
 - Run the gradient descent starting from $x^0 = [-2, -8]$. Does it converge to the global minimum?
 - Plot the log error and estimate the rate of convergence from the logarithmic plot.
 - Try different step sizes. Which step sizes yield a faster convergence? Which are more accurate?